

YASH AGARWAL

Data Engineer | Python · SQL · PySpark · AWS · GCP · ETL Pipelines

yashagarwala2709@gmail.com · +91-8003664759 · [linkedin.com/in/yashagarwal-ai](https://www.linkedin.com/in/yashagarwal-ai) · github.com/yashagarwal16 · Open to Relocation Across India

TECHNICAL SKILLS

Languages: Python · SQL

Data Engineering: ETL Pipelines · Star Schema · Medallion Architecture · Incremental Loads · Data Modeling · Data Validation

Cloud: AWS (S3, Lambda, EC2, RDS) · GCP (Dataproc, BigQuery, GCS)

Big Data: Apache Spark / PySpark · Delta Lake

Databases: SQL Server (SSMS) · MySQL · MongoDB

Tools: Docker · Git / GitHub · CI/CD · Pandas · NumPy

PROFESSIONAL EXPERIENCE

Data & Python Developer Intern · L. R. Tech Solutions

Oct 2025 – Present

Python · Data Analysis · Cloud Data Migration · SQL · Pandas

- Migrated raw business data from on-premise systems to a cloud storage layer, built Python ingestion scripts, validated data integrity post-migration, and restructured data into analysis-ready formats.
- Built and delivered an end-to-end ML model for a client firm — handled data cleaning, feature engineering, model training and evaluation from scratch using Python.
- Conducted data analysis across multiple internal datasets using Python and SQL, producing reports and reusable analytical scripts that reduced ad-hoc analysis time for the team.

Data Pipeline Intern · Linux World Informatics Pvt. Ltd.

Jul – Dec 2024

Python · AWS (EC2, Lambda, S3) · REST APIs · Pandas

- Automated data ingestion from multiple sources (CSV, JSON, REST APIs) using Python, eliminating manual preparation steps and reducing data processing time by 60%
- Deployed data-serving infrastructure on AWS (EC2 / Lambda) with REST APIs handling 500+ daily requests at 99.5% uptime; added structured logging for pipeline monitoring

Cloud & ETL Intern · Edunet Foundation — AICTE × IBM SkillsBuild

Jul – Aug 2025

Python · IBM Cloud · AWS Lambda · Pandas

- Built ETL workflows on IBM Cloud processing 5,000+ records/day using Python; reduced data pipeline latency by 30% through serverless AWS Lambda optimisation

PROJECTS

E-Commerce Lakehouse — Medallion Architecture

2025-2026

PySpark · GCP Dataproc · BigQuery · GCS · Delta Lake

- Built a 3-layer Medallion pipeline (Bronze → Silver → Gold) on GCP Dataproc using PySpark — ingested the Brazilian E-commerce dataset (100K+ orders, 9 tables), applied Silver-layer cleaning and multi-table joins, and computed Gold-layer customer RFM and seller SLA metrics using window functions
- Optimised PySpark jobs using broadcast joins and partition pruning; loaded final Gold tables into BigQuery as the analytics serving layer

SQL Data Warehouse & ETL Pipeline

2025

Python · SQL Server (SSMS) · Docker · Star Schema · CI/CD

- Designed a star schema warehouse (3 fact + 8 dimension tables) with automated Python ETL processing 10K+ daily records at 99.8% accuracy; improved query performance by 45% through indexing and partitioning (8s → 4.4s avg execution)
- Containerised the full pipeline with Docker and configured CI/CD with incremental load scheduling — reducing full-refresh time by 70%

EDUCATION

B.Tech — Artificial Intelligence & Data Science

Expected May 2026

Arya College of Engineering & IT, Jaipur · CGPA: 8.3 / 10.0

CERTIFICATIONS

AI & Cloud — AICTE × IBM SkillsBuild (2025) · MongoDB for Developers — MongoDB University (2024) · Python — GUVI (2023)